

Technical Architecture Document

We propose a **governed execution layer** for agentic AI systems that enforces safety, compliance, and observability by design. This layer wraps around existing agent frameworks and provides a Python-first API while leveraging Rust for performance-critical tasks. The architecture is **modular**: each component (orchestration, memory, logging, policy, etc.) is isolated and extensible. At runtime, a central **Orchestrator** coordinates multiple specialized agents (each a callable module or model) to carry out multi-step workflows. This follows the standard orchestration pattern – “a network of AI agents... each designed for specific tasks, working together to automate complex workflows” ¹. The orchestrator activates the right agents in sequence (or in parallel) and routes intermediate outputs as needed, similar to a “digital symphony” that aligns diverse agent capabilities ².

- **Python-First API, Rust Core:** The developer interface is Pythonic (pip-installable SDK) for ease of use, introspection, and rapid prototyping. In critical paths (e.g. batched model inference, data serialization), we implement components in Rust to avoid Python’s garbage-collector overhead. Benchmarks show identical AI agent workloads use ~5x less RAM in Rust than in Python ³, highlighting Rust’s efficiency for high-throughput agent calls.
- **Integration with Agent Frameworks:** We support plug-ins for popular agent stacks. For example, DSPy (Declarative Self-improving Python) and LangGraph each define graph-based multi-agent workflows ⁴ ⁵. AutoGen (Microsoft’s multi-agent framework) provides event-driven workflows for “deterministic and dynamic agentic workflows for business processes” ⁶. Our orchestrator can import or call these agents/functions internally. In practice, developers can wrap LangGraph or DSPy flows into our layer, inheriting governance without reworking existing code.

Core Components

- **Agent Orchestration:** A central coordinator dispatches subtasks to specialized agents or tools. Each agent is a node in the workflow graph, and the orchestrator manages the edges between them. Agents communicate via structured messages (JSON/YAML states) so that control and data flow are explicit and auditable. This aligns with IBM’s definition of orchestration as coordinating multiple specialized agents “within a unified system” to achieve shared goals ¹. The orchestrator supports sequential and concurrent patterns (e.g. pipeline, branching, feedback loops) and can involve human-in-the-loop (HitL) steps as needed for oversight.
- **Memory Store & Policy:** Agents often need context or memory between steps. We include a secure memory module (e.g. a vector DB or database) to store agent state, action history, and long-term facts. This memory is governed by policies: data is encrypted at rest, and fine-grained retention rules ensure compliance. For example, EU privacy laws require that personal data be deletable on request (“right to be forgotten”) ⁷. Our memory layer supports such data purge policies and redaction. Role-based access controls (RBAC) and audit logs govern read/write access to memory. We enforce data lineage: every memory entry tracks its provenance so that decisions can be traced back to original inputs ⁸ ⁷.

- **Decision Logging (Audit Trail):** Every agent action and output is logged to a write-once audit store before being applied. We record structured entries (e.g. JSON) with fields like timestamp, agent ID, input hash, output summary, confidence, and policy version. This “immutable audit trail” is the foundation of trust: as a best practice, *“every AI call... should produce a structured record before the response hits the user”* ⁹. Logging is append-only (no deletes) and cryptographically signed or checksummed. These logs support forensic review, compliance audits, and rollback. They also feed the observability system (see below) and enable answerable questions like “which model produced that output six months ago” ¹⁰ ¹¹.
- **Policy Engine & Governance Hooks:** A central policy engine enforces high-level rules on agent behavior. Policies are versioned configuration (Policy-as-Code) specifying triggers (e.g. “domain in [healthcare, finance]”, or low confidence) and actions (“flag for review”, “abort”, “require human approval”) ¹². At runtime, each agent output is checked against active policies. Non-compliant outputs are quarantined, sent to human review queues, or blocked entirely. This makes compliance proactive: if an agent requests a tool or produces content, we automatically apply domain-specific constraints (e.g. HIPAA rules if PII detected). Policy violations are logged as part of the audit trail. We also integrate hooks into enterprise controls: for example, before executing a financial transaction, the orchestrator might invoke a bank’s API requiring approval, ensuring external compliance.
- **Evaluation & Rollback Mechanisms:** Agents run continuously, so we build in safety nets. New models or workflow versions are canaried: we run “old” and “new” versions in parallel on sample inputs to compare outputs before full rollout (as advised by governance best practices) ¹⁰. Continuous evaluation metrics (accuracy, bias, relevance) are computed on each output. If anomalies or regressions appear, the orchestrator can roll back to a previous known-good state. All versions of models, code, and policy configurations are tracked in version control. In effect, the system supports “time travel” to a prior configuration in the audit logs, enabling swift rollback and replay of decisions.
- **Compliance Controls (SOC2/GDPR/etc.):** From the ground up, we target enterprise compliance. Data in motion is encrypted (TLS), data at rest is encrypted (AES-256). We support multi-tenant isolation and single sign-on (SSO) integration. Our architecture generates the artifacts needed for audits: full audit logs (for SOC2 “logging and monitoring” requirements), encryption key management, multi-region backups, and access control lists. We incorporate privacy-by-design: for instance, users’ personal data in agent memory is tagged and encrypted separately to enforce GDPR user rights ⁷. As one compliance guide notes, for regulated AI you need **“provable controls, traceability, and accountability at every stage”** ¹³. Our platform provides built-in mechanisms (audit logs, human approval gates, explicit data catalogs) to meet those requirements.
- **Security & Observability:** Beyond compliance, we follow best practices for secure, observable systems. Each component (API server, agents, DBs) emits standard telemetry. We use logging, metrics, and distributed tracing (OpenTelemetry) for traditional observability (CPU/GPU load, request rates, error rates). Crucially, we add AI-specific observability: telemetry on token usage, latency, and output quality. This “fourth pillar” of AI observability is needed to catch issues like hallucinations or model drift ¹⁴ ¹⁵. For example, we track *time-to-first-token* and end-to-end latency per agent; and we flag if a model’s accuracy or safety signals deviate from baselines. Dashboards provide alerts on budget spikes (e.g. if a model suddenly generates far more tokens) and on policy violations detected. In short, we instrument the entire agent workflow so that the operations team can answer “Is this response accurate? Safe? Compliant?” in real time ¹⁵.

- **Scalability:** The execution layer is cloud-native and elastic. Agents are containerized microservices (or lightweight threads) that can scale horizontally. Workflows are stateless except for the governed data stores (memory, logs), so adding nodes increases throughput. We leverage Kubernetes-style orchestration for scaling and resilience (auto-restart, load balancing). For throughput-heavy tasks (e.g. batch retrieval or large-model inference), we can spawn Rust-based service replicas to maximize parallelism. The data plane uses asynchronous, event-driven flows to avoid bottlenecks. This design lets the platform handle both interactive agent calls (low latency) and long-running batch pipelines at scale.
- **Developer Experience:** We provide a Python SDK and CLI so that engineers can integrate agents and governance with minimal friction. A typical developer experience: import our `GovernedAgent` class, plug in a model or tool (even a LangGraph or AutoGen agent), and the system automatically wraps it with governance middleware. Performance-critical internals (batch tokenization, vector search) are in Rust behind the scenes, but the Python interface remains ergonomic. We supply built-in connectors for popular tools (e.g. vector DBs, knowledge bases, custom APIs) so engineers work at a high level. All governance features (logging, policy) are enabled by default, with sensible settings for sandbox or production modes. The goal is: *“what works in development drops into production securely”*, without duplicating compliance work.
- **Enterprise Compliance Readiness:** For certifications like SOC2 or ISO-27001, our platform allows customers to select hosting options: fully-managed cloud, private cloud, or on-premise. Each deployment supports SSO (OAuth/SAML), automated backup and key management, and region-based data control (important for GDPR data residency). We document our security controls and offer transparency (e.g. system architecture reviews). By design, the platform meets the key domains that SOC2 requires: availability, security, integrity, confidentiality, and privacy. We follow NIST guidance and OWASP Top 10 for AI systems. Overall, the architecture embeds compliance hooks so that scaling out the system inherently scales up auditability and security.

Throughout, we emphasize **observability and auditability** at every level. All data exchanges, agent actions, and system events are traceable. This architecture ensures that in high-stakes domains (finance, healthcare, legal, etc.), agentic workflows run in a *known, controlled, and transparent* manner – enabling safe automation without compromising trust.

9 14

Product Vision Document

The Problem

AI agents are rapidly being deployed to make critical decisions, but without proper governance they pose serious safety and compliance risks. These agentic systems operate autonomously and opaquely, leading to unpredictable outcomes in high-stakes domains. For example, an AI agent in healthcare might misdiagnose a patient, or an algorithm in finance might make an unexplainable loan decision, with potentially severe consequences. As industry observers warn, agentic AI autonomy introduces risks (biases or errors) especially in domains like healthcare, finance, or legal services ¹⁶. Regulators and businesses alike recognize this gap: AI deployments now **“demand provable controls, traceability, and accountability”** at

every stage ¹³ . Without a governance layer, organizations cannot “**explain outcomes, trace them back to data and models, and intervene when risks emerge**” ¹⁷ . In short, *ungoverned agentic systems are unsafe for regulated, high-stakes environments*.

Why Now

Several converging trends make this the perfect time for an agentic governance platform:

- **Explosive AI Agent Adoption:** Enterprises are fast embracing multi-agent workflows. By 2025, an estimated 35% of organizations plan to deploy AI agents, rising to 86% by 2027 ¹⁸ . These agents (powered by LLMs and automated tools) promise efficiency but also introduce new complexity. As complexity outpaces human oversight, the need for structured governance becomes urgent.
- **Regulatory Pressure & Compliance:** Governments and regulators are enacting strict AI rules. The EU AI Act (effective 2024) classifies AI risk levels with hefty fines (up to €35M) for violations ¹⁹ . Privacy laws (GDPR, HIPAA) explicitly apply to AI data usage ²⁰ , and enterprises now *expect* SOC2-level security for any AI service ²⁰ ²¹ . In regulated industries (finance, healthcare, government), any AI system must now demonstrate explainability and oversight. For example, banks require AI model risk management and clear audit trails for credit decisions ²² . The bottom line: AI governance is no longer optional. Organizations moving from pilots to production find that **infrastructure and governance gaps** – not model accuracy – are the main barriers to scale ²³ .
- **Demand for Explainability and Trust:** Across industries, customers and regulators are demanding transparency. The global Explainable AI market is booming (projected from ~\$7.8B in 2024 to \$21.1B by 2030 ²⁴). Users need to know *why* a decision was made, especially when lives or large sums are at stake. Without platform-level tools for logging decisions and surfacing reasoning, companies cannot meet these demands. Surveys find that bias, lack of accountability, and “black box” models are among the top concerns slowing AI adoption. Our solution directly addresses this need by embedding traceability and policy enforcement into the core AI workflow.
- **Lack of Existing Infrastructure:** Today’s AI stack is piecemeal. There are frameworks for building agents (LangChain, AutoGen, etc.) but none provide enterprise-grade governance out of the box. Fragmented tools mean teams often bolt on logging or monitoring as an afterthought. As one analysis notes, scaling AI in regulated sectors is hampered by “gaps in infrastructure, governance, and operational maturity” ²³ . We fill this platform gap: no one has yet packaged agent orchestration, compliance and observability into a developer-friendly system. The market is ripe for an integrated solution.

Market Opportunity (TAM/SAM/SOM)

We target a multi-billion-dollar opportunity across several verticals:

- **Legal and Compliance (TAM):** The *legal AI* market alone is growing rapidly – projected from ~\$1.45B in 2024 to ~\$3.90B by 2030 ²⁵ – driven by demand for automated contract review, eDiscovery, and risk analysis. More broadly, enterprise *compliance software* (GRC platforms) is already a ~\$33B market in 2024, expected to exceed \$75B by 2032 ²⁶ . As businesses adopt AI, much of that spending will

shift to intelligent governance tools. Our initial wedge will capture a share of these budgets by servicing corporate legal departments, audit/ compliance teams, and law firms.

- **Sports Analytics – NFL/NIL & European Football (TAM):** Sports teams and leagues are investing heavily in analytics and AI. The global sports analytics market was about \$5.8B in 2025, projected to reach ~\$31B by 2034 ²⁷. European sports analytics (including football/soccer) alone was ~\$1.36B in 2024 and growing at ~19% CAGR ²⁸. In the U.S., name-image-licensing (NIL) valuations and player analytics are emerging use cases for autonomous agents (e.g. contract analysis, performance scouting). This sector will pay premium for explainable insights (e.g. on player value or game strategy). By starting with sports (college/NFL analytics, and top European clubs), we tap a lucrative niche with clear high-stakes decisions and regulatory oversight (e.g. NCAA rules).
- **Other Early Verticals:** We will also serve compliance-intensive domains like F1 financing, insurance underwriting, and government contracts, which exhibit similar needs. These markets add further billions to our SAM. For example, AI in healthcare is projected to be an **\$868 billion opportunity by 2030** ²⁹, and finance/insurance heavily regulate automated risk decisions. As we expand, we can sell to these sectors using the proven pattern of auditability and security.
- **Monopolistic Potential:** By first consolidating governance in these verticals, we establish a strong foothold. The platform becomes “sticky”: once a client has instrumented all their agents with our layer, switching becomes costly (audit trails, compliance evidence, and trained review processes are in our lockstep). We anticipate network effects where industry consortia begin to standardize on our system for AI audits. This path mirrors how Stripe became the default for online payments or Datadog for infrastructure monitoring. We are positioning to be the **de facto standard for agentic decision operations**.

Serviceable Markets and Shares: Initially, our SAM focuses on North American and European enterprises in legal, regulated finance, and professional sports. These regions account for the majority of AI compliance and analytics spend. Even a modest penetration (a few percent of each TAM) yields tens of millions in early revenue (SOM), with rapid growth as AI adoption accelerates.

18 23

Go-to-Market Strategy

Our *wedge* is to onboard high-value “spearhead” customers who face stringent governance needs:

1. **Legal Tech and Corporate Counsel:** We integrate with legal AI tools (e.g. contract analysis bots) to provide audit logs and compliance checks. Early customers include law firms and in-house counsel teams that manage regulatory risk.
2. **Regulatory Compliance Units:** We target banks and insurers with automated decision agents (credit scoring, claims automation). These units already invest in GRC software; adding our layer bridges them into AI.
3. **Sports Teams and Leagues:** Partner with data analytics departments (NFL franchises, sports federations) to govern AI-driven valuation and strategy tools. These organizations have deep pockets and a strong need for explainability (e.g. justifying player contracts or NCAA compliance).

Once established, we will broaden into **healthcare (AI diagnostics pipelines)** and **broader enterprise ops** (automating HR, finance, supply chain decisions) – areas identified by analysts as needing both automation and auditability. In each vertical, we emphasize our turnkey integration with the agent tools they already use (DSPy, LangGraph, AutoGen, in-house RAG systems, etc.), making adoption low-friction.

Differentiation

Our platform is unique in combining these capabilities:

- **End-to-End Governance:** Unlike simple monitoring tools, we enforce policy **inline** with execution. Decisions are *governed*, not just observed after the fact.
- **Traceability and Audit by Default:** Every agent call is logged and versioned. We treat traceability as a feature, not a post-hoc add-on ⁹.
- **Plug-and-Play with Existing Agent Stacks:** We support popular frameworks out-of-the-box (e.g. DSPy, LangGraph, AutoGen). Customers don't have to rewrite; our layer wraps their agents.
- **Policy-Enforcement Engine:** Domain-specific rule engines (e.g. HIPAA compliance for healthcare agents) are built-in. This goes beyond generic observability – it *controls* behavior.
- **Enterprise-Grade Security:** We target the same customer profile as leading SaaS (Stripe, AWS, Datadog) by meeting SOC2/GDPR and sector regulations. This contrasts with developer-only tools that lack hardened controls.

In effect, we position as “*the Stripe for decision infrastructure*” — developer-friendly to integrate, enterprise-grade to trust. No other platform currently offers this combination for agentic AI, which we believe creates strong defensibility.

25 26

Conclusion

Ungoverned AI agents are an existential risk for regulated businesses. With agents now ubiquitous and regulations tightening, a comprehensive governance platform is urgently needed. Our startup is uniquely positioned to capture this opportunity by delivering *observability, control, and compliance at the agentic layer*. We are building the foundational infrastructure for trustworthy AI decision-making — aiming to be the Stripe/Datadog of agentic decision operations.

By focusing initially on high-stakes niches (legal, compliance, sports analytics) and then expanding to healthcare, finance, and beyond, we target a multi-trillion dollar ecosystem. Our deeply technical, policy-aware solution addresses an unmet need, and our early traction with key industry partners will create barriers to entry. We invite investors and partners to join us in defining the future of safe AI automation.

24 27

3 Python Consumes 50x More RAM?! The Rust Benchmark Forcing the AI Industry to Rewrite. | by Prem Chandak | Feb, 2026 | Medium

https://medium.com/@premchandak_11/python-consumes-50x-more-ram-the-rust-benchmark-forcing-the-ai-industry-to-rewrite-6b0d56ec8d1b

4 DSPy

<https://dspy.ai/>

5 Agentic-RAG using AutoGen and Langchain LangGraph Framework | by Shradha Agarwal | Medium

<https://medium.com/@shradhacea/agentic-rag-using-autogen-and-langchain-langgraph-framework-89ac2d684702>

6 AutoGen — AutoGen

<https://microsoft.github.io/autogen/stable//index.html>

7 19 20 21 AI Agent Compliance: GDPR SOC 2 and Beyond | MindStudio

<https://www.mindstudio.ai/blog/ai-agent-compliance/>

8 13 17 22 23 AI Automation Challenges in Regulated Industries | Domino.ai

<https://domino.ai/blog/ai-automation-regulated-industries>

9 10 11 12 Build AI Governance Into Your Solo Dev Stack: A Practical Workshop - DEV Community

https://dev.to/software_mvp-factory/build-ai-governance-into-your-solo-dev-stack-a-practical-workshop-2feg

14 15 What is AI Observability? Key to Monitoring Your LLM Infrastructure | Kong Inc.

<https://konghq.com/blog/learning-center/guide-to-ai-observability>

16 18 Human-in-the-Loop (HitL) Agentic AI for High-Stakes Oversight 2026

<https://onereach.ai/blog/human-in-the-loop-agentic-ai-systems/>

24 Explainable AI Market Size & Share | Industry Report, 2030

<https://www.grandviewresearch.com/industry-analysis/explainable-ai-market-report>

25 Legal AI Market Size, Share & Trends | Industry Report, 2030

<https://www.grandviewresearch.com/industry-analysis/legal-ai-market-report>

26 Compliance Management Software Market Size And Forecast

<https://www.verifiedmarketresearch.com/product/compliance-management-software-market/>

27 Sports Analytics Market Size, Share, Global Growth Report, 2034

<https://www.fortunebusinessinsights.com/sports-analytics-market-102217>

28 Europe Sports Analytics Market Size & Outlook, 2030

<https://www.grandviewresearch.com/horizon/outlook/sports-analytics-market/europe>

29 AI's US\$ 868 billion healthcare revolution | Strategy&

<https://www.strategyand.pwc.com/de/en/industries/pharma-life-sciences/ai-healthcare-revolution.html>