

v1

Absolutely — here's a **full expanded technical document** integrating **Graph-R1** and **Agent-R1** into your hybrid agent memory architecture. This is a detailed design doc that includes pipeline, implementation, and reasoning for why each step improves the system.

Expanded Hybrid Agent Memory Architecture

Structured + Neural + RL (Memory-R1 / Graph-R1 / Agent-R1)

Author: [Your Name]
Date: 2025-12-23

1. Overview

This architecture integrates:

1. **Structured canonical memory** – deterministic, entity-centric facts
2. **Neural long-term memory (Titans/MIRAS)** – learned embeddings with online updates
3. **RL-based memory and reasoning agents** – combining Memory-R1, Graph-R1, and Agent-R1 concepts

Goals:

- Maintain explainable memory and deterministic reasoning
- Enable long-term generalization and adaptive memory
- Efficient multi-step reasoning with learned policies
- Minimize hallucinations and optimize retrieval efficiency

Key Contributions of Integrated Concepts:

Concept	Role
Memory-R1	RL-based memory CRUD decisions (learn ADD/UPDATE/DELETE/NOOP)

Graph-R1 Multi-turn RL reasoning over hypergraph memory, learned path selection

Agent-R1 Unified RL agent framework, formalizes states, actions, rewards for training Memory/Reasoning agents

2. System Pipeline

1. Observation Layer

└ raw inputs + embeddings

2. Canonicalization Layer

└ structured S-P-O extraction → entity-centric graph

3. Candidate Recall + Soft Clustering

└ embedding retrieval + semantic clustering

4. RL Agents (Memory-R1 + Graph-R1 + Agent-R1)

└ Memory Manager Agent (Memory-R1) → decides memory CRUD

└ Graph-R1 Multi-Turn Reasoner → navigates hypergraph, multi-hop reasoning

└ Answer Agent (Memory-R1) → selects best clusters/facts

5. Neural Memory Module (Titans/MIRAS)

└ pattern memory for generalization, online updates using surprise metric

6. Query Resolution

- ↳ deterministic consolidation of facts

7. LLM Projection

- ↳ natural language answer generation

3. Layer Details & Improvements

3.1 Observation Layer

- **Input:** raw text, events, embeddings
- **Improvement:** embeddings feed RL agent states; captures immediate context for downstream decisions
- **Implementation:** SentenceTransformers, tokenized embeddings

3.2 Canonicalization Layer

- **Function:** extract structured (subject, predicate, object) facts, link to graph nodes
- **Improvement:** ensures deterministic memory; reduces LLM hallucination by enforcing explicit facts
- **Implementation:** spaCy/NER + custom S-P-O extractor, timestamped nodes

3.3 Candidate Recall + Soft Clustering

- **Function:** retrieve relevant clusters via embedding similarity; merge similar facts
- **Improvement:** reduces retrieval noise, improves multi-hop reasoning efficiency
- **Implementation:** Agglomerative/KMeans clustering on embeddings

3.4 RL Agents: Memory-R1 + Graph-R1 + Agent-R1

Memory Manager Agent (Memory-R1)

- **State:** new candidate fact + existing graph
- **Actions:** ADD / UPDATE / DELETE / NOOP
- **Reward:** downstream QA correctness
- **Improvement:** transforms memory CRUD into learnable RL problem; reduces manual heuristics
- **Implementation:** PPO or GRPO agent with discrete actions

Graph-R1 Multi-Turn Reasoner

- **State:** current query, graph snapshot, candidate clusters
- **Action:** traverse hypergraph edges, expand nodes, select clusters iteratively
- **Reward:** combined path validity + answer correctness
- **Improvement:** agent learns optimal multi-hop reasoning over hypergraph, replacing brittle traversal heuristics
- **Implementation:** hypergraph stored in NetworkX, RL agent selects edges, multi-turn episodes for training

Answer Agent (Memory-R1)

- **State:** query + retrieved clusters
- **Action:** select subset of clusters/facts for final reasoning
- **Reward:** correctness of final answer
- **Improvement:** filters irrelevant facts, reduces hallucinations, optimizes reasoning efficiency
- **Implementation:** PPO agent on cluster embeddings

Agent-R1 Training Framework

- **Role:** unifies RL agents' states, actions, transitions, rewards
- **Improvement:** ensures standardized, modular, and end-to-end trainable RL setup for Memory + Graph Reasoning
- **Implementation:** define MDP for all RL agents, train with shared reward signals

3.5 Neural Memory Module (Titans/MIRAS)

- **Function:** learned long-term memory embeddings; online updates on surprise
 - **Improvement:** generalizes reasoning to unseen or partially observed facts; complements structured memory
 - **Implementation:** lightweight feedforward or transformer module; gradient updates on surprising inputs
-

3.6 Query Resolution

- **Function:** deterministic consolidation of selected facts
 - **Improvement:** ensures explainable final answer; combines RL-selected clusters with canonical facts
 - **Implementation:** graph traversal or aggregation logic
-

3.7 LLM Projection

- **Function:** verbalizes resolved facts into natural language answer
 - **Improvement:** human-readable interface, optionally integrates neural memory hints
 - **Implementation:** GPT or local LLM
-

4. Implementation Notes

Component	Prototype Implementation
Structured Graph	NetworkX / SQLite
Embeddings	SentenceTransformers / OpenAI embeddings
Soft Clustering	Agglomerative / KMeans
RL Agents	Memory-R1: PPO / GRPO; Graph-R1 multi-turn RL; Agent-R1 MDP framework
Neural Memory	Small feedforward/transformer embeddings
Query & LLM	GPT / local LLM

Training Strategy:

- Reward = final answer correctness + minimal steps + low hallucination
 - Start with 5–10 entities, 10–20 facts, 2–3 query types for prototype
-

5. How Each Step Improves Hybrid Design

Layer	Improvement
Observation	Provides rich embeddings for RL states
Canonicalization	Deterministic structured memory, reduces LLM hallucination
Soft Clustering	Reduces noise, merges similar facts for efficient reasoning
Memory-R1 Memory Manager	Learned memory CRUD, reduces manual heuristics
Graph-R1 Multi-Turn Reasoner	Learns multi-hop reasoning, improves path selection
Answer Agent	Filters relevant clusters, reduces retrieval noise
Agent-R1 Framework	Standardized RL training, modular and scalable
Neural Memory	Long-term generalization, supports unseen facts
Query Resolution	Consolidates deterministic answer from structured memory

6. References

1. Memory-R1: *Learning to Manage Memory for LLMs via RL*, arXiv:2508.19828
 2. Graph-R1: *Agentic Multi-Turn RL over Hypergraph Memory*, GitHub: LHRLAB/Graph-R1
 3. Agent-R1: *End-to-End RL Framework for LLM Agents*, arXiv:2511.14460
 4. Titans + MIRAS: Google Research, 2025
 5. LangChain / Zep: structured graph + embeddings pipelines
-

Summary:

This **expanded hybrid architecture** combines **structured memory**, **neural memory**, and **state-of-the-art RL agents**(Memory-R1, Graph-R1, Agent-R1) to achieve **explainable, adaptive, and scalable memory-driven reasoning** for LLM-based agents. Each layer contributes either **determinism, generalization, or learned reasoning policies**, forming a cohesive end-to-end system.

If you want, I can also **draw a detailed system diagram** showing **all layers, RL agents, neural memory, and data flow**, which would make this a complete design blueprint for implementation.

Do you want me to do that next?

v2

Absolutely — we can extend this **hybrid memory + RL reasoning architecture** to handle **multimodal data** and **complex domains like finance or policy-driven markets**. Let's break it down carefully.

1 Multimodal Memory Integration

Your current hybrid system is mostly text-based, with structured S-P-O graphs, embeddings, and neural memory. To extend to **images, charts, videos, or audio**, we need to:

Observation Layer: Multimodal Inputs

- **Text:** news, reports, policy documents
- **Images:** stock charts, satellite imagery, technical diagrams
- **Audio/Video:** earnings calls, press conferences

Implementation:

- Encode each modality into **vector embeddings**:
 - Text → BERT / OpenAI embeddings
 - Images → CLIP / ViT embeddings
 - Audio → Wav2Vec or similar
 - Include a **modality tag** in embeddings for downstream agents.
-

Canonicalization Layer: Structured Multimodal Facts

- Expand (Subject, Predicate, Object) to (Subject, Predicate, Object, Modality, Timestamp)
- Examples:
 - (Apple, StockTrend, Upward, ImageChart, 2025-12-23)
 - (Fed, InterestRatePolicy, Increase, Text, 2025-12-22)

Improvement:

- Explicitly tracks which modality contributes to each fact
 - Enables cross-modal reasoning, e.g., "Combine news + chart patterns for trend prediction"
-

Candidate Recall & Soft Clustering

- Use **joint multimodal embeddings** for clustering:
 - Combine text + image vectors via a projection layer
 - Soft clustering groups similar patterns across modalities

Benefit:

- Reduces noise across different modalities
 - Allows RL agents to reason over **multimodal clusters**, e.g., news + chart + social sentiment
-

2 RL Agents Adaptation (Memory-R1 / Graph-R1 / Agent-R1)

Memory Manager

- Decisions now consider **modality** when updating memory:
 - Should a news report override a chart pattern?
 - Should older policy documents be deleted or flagged?
- **Reward signal:** downstream prediction accuracy, multi-step financial metrics (profit, portfolio performance, risk)

Graph-R1 Reasoner

- Hypergraph edges now represent **cross-modal relationships**:
 - (FedRateIncrease, affects, S&P500, Text→MarketDataChart)
- Multi-turn reasoning: agent can traverse across text, chart, and audio nodes to make informed predictions.

Answer Agent

- Selects the most relevant multimodal clusters for decision-making:
 - For stock prediction: clusters of news, chart patterns, macroeconomic indicators
 - For policy impact: clusters of regulatory announcements + historical market reactions
-

3 Neural Memory for Multimodal Generalization

- Titans/MIRAS-style embeddings now include **cross-modal projections**:
 - Learn to map text, images, and audio into a **shared latent space**
 - Online updates incorporate new data from any modality
 - Helps **generalize across unseen combinations**:
 - Example: a policy announcement + a correlated chart pattern → predicted market effect
-

4 Query Resolution + LLM Projection

- The final step can now:
 - Integrate multimodal facts into structured reasoning
 - Generate predictions, summaries, or alerts in natural language
 - Include visual outputs: charts, annotated images, graphs
-

5 Domain-Specific Adaptation: Stock Market & Policy

Observation Layer Inputs

- Stock prices & technical indicators
- Market news & social media sentiment
- Government policies, regulatory announcements
- Company reports, earnings call transcripts
- Macroeconomic data (inflation, interest rates)

Memory / Graph Representation

- Nodes:
 - Companies, indices, sectors, regulators, macro variables
- Hyperedges:
 - Market impact relationships: **Company** → **Sector** → **Index**

- Causal/temporal connections: FedRateDecision → MarketTrend

RL Agent Objectives

- Memory Manager:
 - Maintain **relevant and accurate financial facts**
 - Discard stale or irrelevant data
- Graph Reasoner:
 - Traverse market-relevant hypergraph for **multi-hop causal reasoning**
 - Example: “Fed policy → interest rate → tech sector → Apple stock”
- Answer Agent:
 - Select clusters for prediction / trading decisions
 - Reward: downstream financial metrics (profit, loss avoidance, portfolio volatility reduction)

Neural Memory Role

- Captures **latent correlations**:
 - Example: “Policy announcement + previous similar earnings report → likely stock move”
 - Helps generalize to **novel market situations**
-

6 Implementation Notes for Multimodal Finance Prototype

Component	Implementation
Graph Memory	Hypergraph with nodes for companies, policies, indices; NetworkX or Neo4j
Embeddings	Text: BERT/OpenAI; Images: CLIP; Audio: Wav2Vec; joint multimodal projection layer
Soft Clustering	Multimodal embeddings + Agglomerative/KMeans
RL Agents	Memory-R1 + Graph-R1 + Agent-R1; PPO/GRPO; rewards based on market prediction accuracy
Neural Memory	Cross-modal embedding module; online updates

Query Resolution	Deterministic hypergraph traversal + weighted aggregation of selected clusters
LLM Projection	Generate predictions, explanations, alerts; optional charts/visuals

7 Benefits of This Extension

- **Cross-modal reasoning:** Text + chart + audio combined into a single decision-making process
 - **Dynamic adaptation:** RL agents learn which modalities are most predictive for a given task
 - **Explainable reasoning:** Hypergraph + canonical facts maintain auditability
 - **Long-term generalization:** Neural memory learns latent patterns across modalities
 - **Domain-specific optimization:** Financial predictions, policy impact, market simulations
-

✓ Summary:

By combining **multimodal embeddings, structured hypergraph memory, RL reasoning (Memory-R1, Graph-R1, Agent-R1), and neural generalization**, your hybrid system can handle **complex, dynamic, multi-source domains** like stock markets, macro policy analysis, and predictive forecasting — all while remaining **explainable, adaptive, and scalable**.

If you want, I can now **draw a complete multimodal hybrid system diagram** showing **cross-modal embeddings, hypergraph memory, RL agents, and neural memory interactions**, tailored for **finance and policy decisions**. This would serve as a visual blueprint for implementation.

Do you want me to do that next?

v3

Self-Improving Agent Memory & Decision System (Technical Design)

1. Purpose and Design Goals


This document specifies an **agent memory and decision-making architecture** that improves over time while remaining:


- **Auditable** (traceable reasoning and decisions)
- **Scalable** (avoids $O(N)$ memory traversal)
- **Multimodal** (text, images, time-series, structured data)
- **Self-improving** (learns what to remember, how to reason, and how to decide)
- **Human-in-the-loop compatible** (explicit feedback and overrides)


The design addresses known failures of vector-only RAG systems and unstructured memory chains by combining:

- Entity-centric symbolic memory
- Soft-clustered neural memory
- Graph-based reasoning
- RL-style policy optimization for reasoning and decisions

2. High-Level Architecture







3. Memory Layer Design

3.1 Symbolic Entity Graph

Purpose: Deterministic state, relational queries, auditability.

- Nodes: canonical entities (Person, Company, Policy, Asset, Event)
- Edges: typed predicates with timestamps

Example:

(FED) --[raises_rate {t=2025Q1}]-> (Interest_Rates)

(Interest_Rates) --[affects]-> (Bank_Stocks)

-

Properties:

- $O(\log N)$ or bounded traversal
- Supports exact queries and filters
- Human-readable and debuggable

3.2 Neural Soft-Cluster Memory

Purpose: Semantic generalization and fuzzy recall.

- Embeddings stored in **overlapping soft clusters** (mixture-of-centroids)
- Each memory belongs to multiple clusters with learned weights

Benefits:

- Avoids brittle hard clustering
 - Enables regime-level recall ("rate-hike cycles", "risk-off periods")
 - Improves candidate recall without full vector scans
-

3.3 Temporal / Episodic Memory

- Stores full interaction episodes
 - Indexed by time, task, and outcome
 - Used for:
 - post-hoc analysis
 - self-improvement
 - human review
-

4. Candidate Recall & Reasoning

4.1 Candidate Recall (Critical Step)

Candidate recall narrows the search space before LLM reasoning.

Sources:

1. Entity graph neighborhood
2. Relevant soft clusters
3. Recent episodic memory

This converts an intractable global memory into a **bounded, high-recall candidate set**.

4.2 Graph-Based Reasoning (Graph-R1 Inspired)

- Reasoning = graph path selection
- Each step conditioned on state and goal
- Paths scored by learned policy

Example:

Policy → Macro → Sector → Asset

This prevents the LLM from hallucinating relationships.

5. Decision Layer

5.1 Multiple Reasoning Policies

For a given task, the system executes multiple strategies:

- News-first
- Market-data-first
- Historical-analogy-first

Each produces:

- Decision
 - Confidence
 - Explanation trace
-

5.2 Decision Selection

A critic selects the final output based on:

- Historical performance
 - Risk constraints
 - Explanation consistency
-

6. Context Trace & Explanation Storage

Every decision produces a **Context Trace Object**:

```
{  
  task_id,  
  retrieved_memories,
```

```
graph_paths_traversed,  
modalities_used,  
reasoning_policy,  
intermediate conclusions,  
final decision,  
confidence,  
timestamp  
}
```

Stored in episodic memory.

Why this matters:

- Full auditability
- Post-mortem analysis
- Human trust

7. Human-in-the-Loop Interaction

Humans can:

- Approve / reject decisions
- Edit explanations
- Provide counterfactuals

Human feedback is stored as:

(feedback_type, target_trace_id, annotation)

This feedback becomes a training signal.

8. Self-Improvement Loops

8.1 Fast Loop – Reasoning Optimization

- Optimizes prompts, graph traversal heuristics
 - Uses critic scores
 - No model retraining
-

8.2 Mid Loop – Policy Learning (RL-style)

- Policies decide:
 - which memories to retrieve
 - which graph paths to explore
 - which modality to trust

Rewards:

- Decision correctness
 - Risk-adjusted outcomes
 - Human approval
-

8.3 Slow Loop – Representation Adaptation

- Updates cluster centroids
 - Adjusts embedding spaces
 - Learns regime shifts
-

9. Multimodal Extension

- Images → vision encoders → event embeddings
- Time-series → pattern embeddings
- Text → entity + semantic embeddings

All modalities map to:

- entity graph (symbolic)
 - soft clusters (neural)
-

10. Failure Modes & Safeguards

Risk	Mitigation
Feedback loops	Ensemble critics + human veto
Overfitting policies	Entropy regularization
Memory bloat	Learned retention + decay
Explanation drift	Trace consistency checks

11. Why This Improves Over Vector-RAG

- No $O(N)$ traversal
 - No black-box memory
 - Explicit reasoning paths
 - Continuous improvement
 - Human-aligned decision making
-

12. Summary

This architecture transforms an LLM from a passive reasoner into a **self-improving, memory-grounded decision system** capable of operating in high-stakes, multimodal, and evolving domains while remaining transparent and controllable.

I've created a **new, expanded technical design document** in the canvas that:

- Improves on the previous hybrid memory design
- Adds **explicit self-improving decision-making layers**
- Introduces **context trace objects** (full reasoning provenance)
- Treats **human interaction and feedback as first-class training signals**
- Keeps the system auditable, scalable, and RL-compatible without black-boxing

What's new and important in this version

1. Context traces are now a core artifact

- Every decision stores *how* it was made:
 - retrieved memories
 - graph paths
 - modalities used
 - reasoning policy
 - intermediate conclusions
- This enables:
 - audits
 - post-mortems
 - human correction
 - offline learning

2. Self-improvement is structured into 3 loops

- Fast loop: reasoning & retrieval optimization (DSPy / OpenPipe-style)
- Mid loop: RL-style policy learning over memory, graph traversal, modality trust
- Slow loop: representation & regime adaptation (Titans / MIRAS flavor)

3. Human-in-the-loop is no longer an afterthought

- Feedback attaches directly to decision traces
- Human edits and vetoes become training data
- Enables alignment without retraining base models

4. Multimodal + market decision-ready

- Images, charts, text, policies all map into:
 - symbolic entity graph
 - neural soft clusters
- Designed for domains like finance, policy forecasting, compliance

5. Clear separation of concerns

- Symbolic memory = truth & structure

- Neural memory = generalization & recall
 - RL policies = *how* the system reasons and decides
-

If you want, next we can:

- Turn this into a **concrete weekend MVP plan** (components + minimal stack)
- Add **pseudo-code for the self-improvement loops**
- Design a **market-specific reward function**
- Stress-test the design against adversarial / regime-shift scenarios

Just tell me where you want to go next.

mvp/prd

Product Requirements Document (PRD)

Self-Improving Agent Memory & Decision System – Weekend MVP

1. Product Overview

Product Name (working): TraceMind

Goal:

Build a minimal but end-to-end self-improving agent memory and decision system that:

- Makes structured decisions (e.g., market / policy / analysis tasks)
- Stores *how* decisions were made (context traces)
- Improves retrieval and reasoning policies over time
- Supports human feedback

This MVP is explicitly scoped to be built over a single weekend, prioritizing architectural correctness over scale.

2. Target Users

Primary

- AI engineers / researchers building agentic systems

Secondary

- Quant / policy analysts experimenting with AI-assisted reasoning
-

3. Core Use Case (MVP)

Example Task:

"Given recent news and market data, should we be bullish, bearish, or neutral on a stock/sector?"

System must:

1. Ingest new information
 2. Store it in structured memory
 3. Retrieve relevant context
 4. Produce a decision + explanation
 5. Store a full reasoning trace
 6. Accept human feedback
 7. Improve next time
-

4. Non-Goals (Explicitly Out of Scope)

- Training or fine-tuning foundation models
 - High-frequency trading or real money execution
 - Full multimodal ingestion beyond basic images/time-series
 - Perfect prediction accuracy
-

5. Functional Requirements

5.1 Ingestion & Canonicalization

Requirements:

- Accept text input (news, user prompts)
- Extract entities and relations
- Generate (Subject, Predicate, Object, Timestamp)

MVP Implementation:

- LLM-based extractor (single prompt)
 - Simple entity normalization (string-based)
-

5.2 Memory Layer

A. Symbolic Entity Graph

Requirements:

- Store entities and typed edges
- Support neighborhood queries

MVP Implementation:

- NetworkX or Neo4j
 - Nodes = entities
 - Edges = predicates with metadata
-

B. Neural Soft-Cluster Memory

Requirements:

- Store semantic memories
- Support fuzzy recall without full vector scan

MVP Implementation:

- Sentence-transformer embeddings
 - K-means centroids (small K)
 - Soft assignment via cosine similarity
-

C. Episodic Memory

Requirements:

- Store full decision traces

MVP Implementation:

- JSONL or SQLite
-

5.3 Candidate Recall

Requirements:

- Retrieve a bounded set of relevant memories

MVP Implementation:

- Graph neighborhood expansion
 - Top-N soft clusters
 - Recent episodic context
-

5.4 Reasoning Layer

Requirements:

- Execute multiple reasoning strategies
- Prevent hallucinated relationships

MVP Implementation:

- Hard-coded strategies:
 - Graph-first
 - Semantic-first
 - Each produces a decision + explanation
-

5.5 Decision Selection

Requirements:

- Choose best reasoning output

MVP Implementation:

- **Rule-based critic:**
 - **explanation length**
 - **memory coverage**
 - **confidence heuristics**
-

5.6 Context Trace Storage (Critical)

Requirements:

Every decision must store:

- **Retrieved memories**
- **Graph paths**
- **Strategy used**
- **Final decision**
- **Confidence**

MVP Implementation:

- **Structured JSON object**
-

5.7 Human Feedback

Requirements:

- **Accept approve/reject**
- **Accept text feedback**

MVP Implementation:

- **CLI or notebook input**
 - **Feedback linked to trace_id**
-

5.8 Self-Improvement (MVP Scope)

Requirements:

- **System behavior changes based on outcomes**

MVP Implementation:

- **Strategy success counters**
 - **Retrieval weighting updates**
 - **No gradient-based RL yet**
-

6. Weekend Build Plan

Day 1 – Core Memory & Reasoning

Morning:

- **Define schemas (Entity, Edge, Trace)**
- **Set up graph store**

Afternoon:

- **Build ingestion + canonicalization**
- **Implement soft-cluster memory**

Evening:

- **Candidate recall pipeline**
 - **First end-to-end query**
-

Day 2 – Decisions, Traces & Learning

Morning:

- **Multiple reasoning strategies**
- **Decision critic**

Afternoon:

- Context trace persistence
- Human feedback loop

Evening:

- Self-improvement heuristics
 - Failure analysis + cleanup
-

7. Success Metrics (MVP)

- Decisions are reproducible from traces
 - Memory retrieval is bounded and explainable
 - System behavior changes after feedback
 - No $O(N)$ memory traversal
-

8. Risks & Mitigations

Risk	Mitigation
Over-comple xity	Hard scope limits
Hallucination s	Graph grounding
Feedback noise	Manual review

9. Tech Stack (Suggested)

- **Python**
 - **NetworkX**
 - **SentenceTransformers**
 - **SQLite / JSONL**
 - **OpenAI / local LLM**
-

10. What This MVP Proves

- **Hybrid memory > vector-only RAG**
- **Traces enable trust & learning**
- **Self-improvement doesn't require retraining**

This MVP serves as a foundation for scaling into a production-grade agent memory and decision platform.

How this PRD is intentionally structured

This is **not** a startup-style fluffy PRD. It's an *engineering-first* PRD designed to:

- Prove the **hybrid memory thesis**
- Validate **trace-based reasoning**
- Demonstrate **self-improvement without retraining**
- Keep scope brutally tight

What makes this MVP realistic for a weekend

- No training loops
- No distributed systems
- No fancy RL yet (policy updates via counters/heuristics)
- Everything inspectable via JSON + graph visualizations

You'll end the weekend with:

- A working agent that **makes decisions**
- A **memory system that doesn't collapse under scale**
- A **trace log that explains every decision**
- A primitive but real **learning loop**

The key thing to internalize

This MVP is *already* doing something most “agent frameworks” don’t:

Learning what *not* to retrieve, trust, or reason over.

That’s the seed of real intelligence.

v4

Self-Improving Agent Memory & Decision System (AgentMem v2)

1. Motivation & Design Principles

Modern agent memory systems fail under scale because they:

- Treat memory as unstructured text blobs (Vector-RAG)
- Conflate retrieval with reasoning
- Lack auditability, causality, and learning signals

This design explicitly separates:

- Memory *representation* (deterministic, structured)
- Memory *control* (learned, policy-driven)
- Reasoning & decision-making (traceable, improvable)

The result is a **hybrid system** combining:

- Entity-centric graph memory
 - Soft semantic clustering
 - Agent-controlled memory operations (AgentMem)
 - Self-improving decision policies using lightweight RL-style feedback (GRPO / bandits)
-

2. High-Level Architecture

[Multimodal Inputs]

|

v

[Canonicalization Layer]

|

v

[Structured Memory Layer] <----> [Soft Cluster Index]

|


v

[Candidate Recall Engine]

|

v

[Reasoning + Planning Layer]



3. Canonicalization Layer

Purpose

Transforms raw inputs into **normalized, machine-operable facts**.

Inputs

- Text (conversation, documents)
- Images / video (via perception models)
- External signals (market data, events, policies)

Outputs

- Entity references
- Predicate-relation triplets
- Temporal & source metadata

Example:

"User says Tesla stock will rally after earnings"
→ (Tesla, expected_movement, rally)
→ (Tesla, catalyst, earnings)
→ confidence=0.42, source=user

This layer acts as a **denoising autoencoder** for memory.

4. Structured Memory Layer

4.1 Entity-Centric Graph Memory

- Nodes: entities (people, assets, concepts)
- Edges: typed, timestamped relations
- Properties: confidence, decay, provenance

Complexity:

- Updates: $O(\log N)$
- Traversals: $O(\log N)$ to $O(k \text{ hops})$

4.2 Episodic Memory Store

Stores **raw interaction traces**:

- Full prompts
- Retrieved memory IDs
- Intermediate reasoning summaries
- Final decisions

Used for:

- Auditing
 - Learning
 - Human review
-

5. Soft Semantic Clustering Layer

Purpose

Improve candidate recall without losing structure.

Mechanism

- Entity and relation embeddings
- Online soft clustering (HDBSCAN / incremental k-means)
- Overlapping membership allowed

Key Property:

Clusters **augment**, not replace, the graph.

6. Candidate Recall Engine

Function

Produces a **bounded, high-recall candidate set**.

Inputs

- Query intent
- Current task context
- AgentMem policy parameters

Retrieval Mix

- Graph traversal (deterministic)
- Cluster expansion (probabilistic)
- Episodic windowing (temporal)

This guarantees **$O(\log N)$** recall with tunable breadth.

7. Reasoning & Planning Layer

Responsibilities

- Compose retrieved facts
- Generate plans or predictions
- Propose actions

Constraints

- Operates only on retrieved memory IDs
- Cannot hallucinate new memory writes

Output

- Decision candidates
 - Confidence estimates
 - Justification summaries
-

8. Decision & Action Layer

Produces a **committed decision**:

- Action taken (or recommendation)
- Confidence
- Risk flags

Example (Markets):

Decision: Reduce exposure to TSLA

Confidence: 0.67

Rationale IDs: [E123, R88, EP45]

9. Context Trace & Provenance System

Every decision stores a **Context Trace**:

Trace Schema

- Input snapshot
- Retrieved memory IDs
- Reasoning summary
- Human inputs (if any)
- Final action
- Observed outcome

This enables:

- Full decision replay
 - Human trust & compliance
 - Offline learning
-

10. AgentMem Controller (Core Innovation)

Role

AgentMem is a **policy over memory operations**, not memory itself.

Action Space

- store(memory)
- link(entity_a, entity_b)
- retrieve(depth, breadth)
- forget(weight)
- defer (do nothing)

State

- Task type
 - Memory growth metrics
 - Recent retrieval effectiveness
 - Human feedback signals
-

11. Self-Improvement & Learning Loop

Phase 1: Heuristics

- Confidence thresholds
- TTL decay
- Fixed retrieval budgets

Phase 2: Bandit Learning

- Reward: decision quality
- Optimize retrieval breadth & storage selectivity

Phase 3: GRPO / RLVR-style Updates

- Offline optimization over traces
- Reward shaping from outcomes + human feedback
- No end-to-end gradient through memory

This avoids instability while enabling learning.

12. Human-in-the-Loop Integration

Humans can:

- Approve / reject decisions
- Correct memory facts
- Provide reward signals

All interactions are stored as **first-class memory objects**.

13. Multimodal Extension

- Images → entities + attributes
- Video → events + timelines
- Charts → structured signals

Memory remains symbolic; embeddings are auxiliary.

14. Failure Modes & Mitigations

Risk	Mitigation
Memory bloat	AgentMem storage policy
Retrieval noise	Candidate recall bounds
Feedback loops	Human-gated rewards
Policy collapse	Offline-only RL updates

15. Why This Works at Scale

- Deterministic core
- Learned control at edges
- Auditable decisions
- Incremental learning

This system scales because **learning controls memory, but structure controls learning**.

16. MVP Implementation Notes

- Graph: Neo4j / SQLite + adjacency lists
- Clusters: FAISS + HDBSCAN
- AgentMem: simple bandit policy
- Traces: append-only log

Weekend-feasible, production-extendable.

prod 2.0

PRD v2 — Self-Improving Agent Memory & Decision System (AgentMem)

1. Overview

Product Name

AgentMem — Self-Improving Agent Memory & Decision System

Problem Statement

Current agent memory systems (Vector-RAG, chain-based updates, monolithic embeddings) fail to scale due to:

- $O(N)$ memory traversal and retrieval
- Semantic interference and hallucinations
- Lack of auditability and causality
- No principled self-improvement loop

AgentMem addresses this by combining **structured memory**, **bounded retrieval**, and **policy-driven self-improvement**, while maintaining full traceability of decisions and human interactions.

2. Goals

G1. Scalable Long-Term Memory

- Support millions of facts without retrieval degradation
- Guarantee bounded retrieval complexity
- Avoid context bloat in LLM inference

G2. Decision-Centric Intelligence

- Store not just facts, but decisions, rationales, and outcomes
- Enable replayable and explainable decision traces

G3. Self-Improvement Without Instability

- Improve retrieval, storage, and decision quality over time
- Avoid end-to-end RL through memory or live inference

G4. Human Trust & Control

- Enable humans to inspect, correct, and guide memory
- Preserve provenance and causality

G5. Multimodal & Real-World Readiness

- Support text, images, time-series, and events
 - Handle domains like markets, policy analysis, and planning
-

3. Non-Goals

NG1. End-to-End Differentiable Memory

- No neural memory that replaces symbolic structure
- No gradient flow through the memory store

NG2. Fully Autonomous Agents

- Human oversight is required for high-impact decisions
- No unsupervised self-modification

NG3. General World Models

- System focuses on task- and entity-centric reasoning
- Not attempting full environment simulation

NG4. Real-Time RL at Inference

- Learning occurs offline or asynchronously
 - Inference-time behavior remains deterministic + bounded
-

4. Target Users

- AI / ML Engineers building agent systems
- Research teams exploring agent memory & reasoning

- Decision-support systems (markets, compliance, ops)
 - Internal tooling requiring explainability
-

5. Functional Requirements

FR1. Canonicalization

- Convert raw multimodal inputs into structured entities and relations
- Attach confidence, timestamps, and sources

FR2. Structured Memory

- Entity-centric graph with typed edges
- Episodic memory for raw traces

FR3. Candidate Recall Engine

- Bounded recall via graph traversal + soft clusters
- Tunable depth and breadth

FR4. AgentMem Controller

- Policy over memory actions (store, retrieve, forget, link)
- State-aware and reward-driven

FR5. Context Trace Storage

- Persist full decision traces
- Enable replay and audit

FR6. Human-in-the-Loop

- Accept corrections, approvals, and feedback
 - Store human inputs as memory objects
-

6. Non-Functional Requirements

NFR1. Performance

- Retrieval latency < 100ms for 99p
- Memory growth linear in facts, not interactions

NFR2. Reliability

- Deterministic retrieval for identical queries
- No silent memory corruption

NFR3. Observability

- Full logging of memory operations
- Metrics for retrieval quality and memory health

NFR4. Safety

- Bounded action space for AgentMem
 - Human gating for irreversible actions
-

7. Success Metrics

Memory Metrics

- Avg. retrieval size per query
- Memory growth rate vs. signal retention
- Graph connectivity entropy

Decision Metrics

- Outcome success rate
- Calibration (confidence vs. accuracy)
- Regret vs. baseline heuristics

Learning Metrics

- Improvement in reward over time
- Reduction in irrelevant retrievals
- Stability of AgentMem policy

Human Metrics

- Correction rate
- Trust / approval ratio

- Time-to-audit per decision
-

8. Risks & Mitigations

R1. Memory Bloat

Risk: Excessive fact storage

Mitigation: AgentMem storage thresholds, decay, human pruning

R2. Retrieval Noise

Risk: Over-broad candidate recall

Mitigation: Hard recall budgets, reward shaping

R3. Feedback Loops

Risk: Self-reinforcing incorrect beliefs

Mitigation: Human validation, offline-only learning

R4. Policy Collapse

Risk: AgentMem converges to trivial behavior

Mitigation: Entropy regularization, policy constraints

R5. Human Overhead

Risk: Excessive manual intervention

Mitigation: Escalation thresholds, selective review

9. Rollout Plan

Phase 1 — MVP (Weekend Build)

- Graph memory + episodic logs
- Heuristic AgentMem
- Manual evaluation

Phase 2 — Learning

- Bandit-based AgentMem
- Reward from outcomes + human feedback

Phase 3 — Advanced Optimization

- GRPO-style offline updates
 - Cross-task policy generalization
-

10. Out of Scope (Explicit)

- Fully autonomous trading
 - Real-time policy learning
 - Black-box neural memory
-

11. Summary

AgentMem v2 is designed to be:

- **Scalable** (bounded memory + retrieval)
- **Auditable** (context traces, provenance)
- **Self-improving** (policy learning without instability)
- **Human-aligned** (oversight by design)

This PRD prioritizes **correctness, trust, and learning efficiency over raw agent autonomy**.

prd 3.0

PRD v2 — Self-Improving Agent Memory & Decision System (AgentMem)

1. Overview

Product Name

AgentMem — Self-Improving Agent Memory & Decision System

Problem Statement

Current agent memory systems (Vector-RAG, chain-based updates, monolithic embeddings) fail to scale due to:

- $O(N)$ memory traversal and retrieval
- Semantic interference and hallucinations
- Lack of auditability and causality
- No principled self-improvement loop

AgentMem addresses this by combining **structured memory**, **bounded retrieval**, and **policy-driven self-improvement**, while maintaining full traceability of decisions and human interactions.

2. Goals

G1. Scalable Long-Term Memory

- Support millions of facts without retrieval degradation
- Guarantee bounded retrieval complexity
- Avoid context bloat in LLM inference

G2. Decision-Centric Intelligence

- Store not just facts, but decisions, rationales, and outcomes
- Enable replayable and explainable decision traces

G3. Self-Improvement Without Instability

- Improve retrieval, storage, and decision quality over time
- Avoid end-to-end RL through memory or live inference

G4. Human Trust & Control

- Enable humans to inspect, correct, and guide memory
- Preserve provenance and causality

G5. Multimodal & Real-World Readiness

- Support text, images, time-series, and events
 - Handle domains like markets, policy analysis, and planning
-

3. Non-Goals

NG1. End-to-End Differentiable Memory

- No neural memory that replaces symbolic structure
- No gradient flow through the memory store

NG2. Fully Autonomous Agents

- Human oversight is required for high-impact decisions
- No unsupervised self-modification

NG3. General World Models

- System focuses on task- and entity-centric reasoning
- Not attempting full environment simulation

NG4. Real-Time RL at Inference

- Learning occurs offline or asynchronously
 - Inference-time behavior remains deterministic + bounded
-

4. Target Users

- AI / ML Engineers building agent systems
- Research teams exploring agent memory & reasoning

- Decision-support systems (markets, compliance, ops)
 - Internal tooling requiring explainability
-

5. Functional Requirements

FR1. Canonicalization

- Convert raw multimodal inputs into structured entities and relations
- Attach confidence, timestamps, and sources

FR2. Structured Memory

- Entity-centric graph with typed edges
- Episodic memory for raw traces

FR3. Candidate Recall Engine

- Bounded recall via graph traversal + soft clusters
- Tunable depth and breadth

FR4. AgentMem Controller

- Policy over memory actions (store, retrieve, forget, link)
- State-aware and reward-driven

FR5. Context Trace Storage

- Persist full decision traces
- Enable replay and audit

FR6. Human-in-the-Loop

- Accept corrections, approvals, and feedback
 - Store human inputs as memory objects
-

6. Non-Functional Requirements

NFR1. Performance

- Retrieval latency < 100ms for 99p
- Memory growth linear in facts, not interactions

NFR2. Reliability

- Deterministic retrieval for identical queries
- No silent memory corruption

NFR3. Observability

- Full logging of memory operations
- Metrics for retrieval quality and memory health

NFR4. Safety

- Bounded action space for AgentMem
 - Human gating for irreversible actions
-

7. Success Metrics

Memory Metrics

- Avg. retrieval size per query
- Memory growth rate vs. signal retention
- Graph connectivity entropy

Decision Metrics

- Outcome success rate
- Calibration (confidence vs. accuracy)
- Regret vs. baseline heuristics

Learning Metrics

- Improvement in reward over time
- Reduction in irrelevant retrievals
- Stability of AgentMem policy

Human Metrics

- Correction rate
- Trust / approval ratio

- Time-to-audit per decision
-

8. Risks & Mitigations

R1. Memory Bloat

Risk: Excessive fact storage

Mitigation: AgentMem storage thresholds, decay, human pruning

R2. Retrieval Noise

Risk: Over-broad candidate recall

Mitigation: Hard recall budgets, reward shaping

R3. Feedback Loops

Risk: Self-reinforcing incorrect beliefs

Mitigation: Human validation, offline-only learning

R4. Policy Collapse

Risk: AgentMem converges to trivial behavior

Mitigation: Entropy regularization, policy constraints

R5. Human Overhead

Risk: Excessive manual intervention

Mitigation: Escalation thresholds, selective review

9. Rollout Plan

Phase 1 — MVP (Weekend Build)

- Graph memory + episodic logs
- Heuristic AgentMem
- Manual evaluation

Phase 2 — Learning

- Bandit-based AgentMem
- Reward from outcomes + human feedback

Phase 3 — Advanced Optimization

- GRPO-style offline updates
 - Cross-task policy generalization
-

10. Out of Scope (Explicit)

- Fully autonomous trading
 - Real-time policy learning
 - Black-box neural memory
-

11. Summary

AgentMem v2 is designed to be:

- **Scalable** (bounded memory + retrieval)
- **Auditable** (context traces, provenance)
- **Self-improving** (policy learning without instability)
- **Human-aligned** (oversight by design)

This PRD prioritizes **correctness, trust, and learning efficiency over raw agent autonomy**.

12. Technical Architecture Overview

AgentMem is composed of four strictly separated layers:

1. **Semantic State & Canonicalization Layer**
2. **Structured Memory Substrate**
3. **Decision & Reasoning Layer**
4. **AgentMem Control & Learning Layer**

This separation ensures learning never corrupts memory, and memory never destabilizes learning.

13. Semantic State & Canonicalization Layer

Purpose

Convert raw multimodal inputs into **low-entropy semantic state** prior to any memory write.

Design (VL-JEPA-Inspired)

- Predicts *task-relevant meaning*, not tokens
- Joint latent space across text, images, charts, and events
- Acts as a semantic denoising autoencoder

Output

- Semantic state vector
- Entity candidates
- Predicate–object hypotheses
- Confidence + provenance metadata

This layer dramatically reduces memory pollution and paraphrase noise.

14. Structured Memory Substrate

14.1 Entity-Centric Graph Memory

- Nodes: entities (people, assets, concepts)
- Edges: typed relations with timestamps, confidence, decay
- Complexity: $O(\log N)$ updates, bounded k-hop traversals

14.2 Episodic Trace Store

Stores raw interaction traces:

- Inputs
- Retrieved memory IDs
- Reasoning summaries
- Decisions
- Human interactions

Used for auditing, replay, and offline learning.

14.3 Decision Memory

Every committed decision is stored with:

- Action taken
 - Supporting memory IDs
 - Confidence
 - Observed outcome (when available)
-

15. Candidate Recall Engine

Function

Produce a bounded, high-recall candidate set for reasoning.

Retrieval Mix

- Deterministic graph traversal
- Soft semantic cluster expansion
- Temporal filtering

Hard limits enforce $O(\log N)$ retrieval complexity.

16. Reasoning & Decision Layer

- Operates only on retrieved memory IDs
- Produces decision candidates, confidence, and rationale summaries
- Cannot write to memory directly

This guarantees separation between reasoning and persistence.

17. Context Trace & Provenance System

Every decision generates a **Context Trace**:

- Input snapshot
- Semantic state summary
- Retrieved memory IDs
- Reasoning summary

- Human feedback
- Final action
- Outcome

Context traces enable explainability, trust, and learning.

18. AgentMem Controller (Self-Improvement Core)

Role

A learned policy that controls *how memory is used*, not what memory contains.

Action Space

- store / suppress
- retrieve(depth, breadth)
- link entities
- decay / forget
- defer

State Signals

- Task type
- Retrieval effectiveness
- Memory growth metrics
- Human feedback

AgentMem actions are schema-constrained and safe by construction.

19. Learning & Optimization Loop

Phase 1: Heuristics

- Static thresholds
- TTL decay
- Fixed retrieval budgets

Phase 2: Bandits

- Reward = outcome quality – retrieval cost
- Learns retrieval breadth and storage selectivity

Phase 3: Offline GRPO / RLVR

- Optimized over stored context traces
 - No gradients through memory or live inference
 - Stable and auditable learning
-

20. Engineering Implementation Notes

Storage

- Graph: Neo4j / SQLite adjacency lists
- Traces: append-only log (Parquet / JSONL)

Indexing

- Embeddings: FAISS
- Clustering: HDBSCAN / incremental k-means

Services

- Memory service (CRUD + traversal)
 - AgentMem policy service
 - Trace logging service
-

21. Key Engineering Invariants

- Memory is deterministic
- Learning is offline
- Retrieval is bounded
- Decisions are replayable

Violating any invariant is considered a system bug.

22. Final Principle

**Semantic compression before storage. Deterministic structure after storage.
Learned control at the boundaries.**